

Lecture 8:

Linking and Loading

Prof. Matt Welsh
September 29, 2009



Topics for today

- Getting from C programs to running executables
- Linking:
 - Process of combining multiple code modules into a runnable program
 - Example: Using standard library code (printf, malloc, etc.) in your own program
- Loading:
 - Process of getting an executable running on the machine
 - Requires resolving references to external symbols in the program (**dynamic linking**)

Example program with two C files

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

void swap()
{
    int temp;

    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
}
```

Example program with two C files

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

Definition of buf[]

Definition of main()

swap.c

```
extern int buf[];

void swap()
{
    int temp;

    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
}
```

Definition of swap()

Declaration of buf[]

Example program with two C files

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

Reference to swap()

swap.c

```
extern int buf[];

void swap()
{
    int temp;

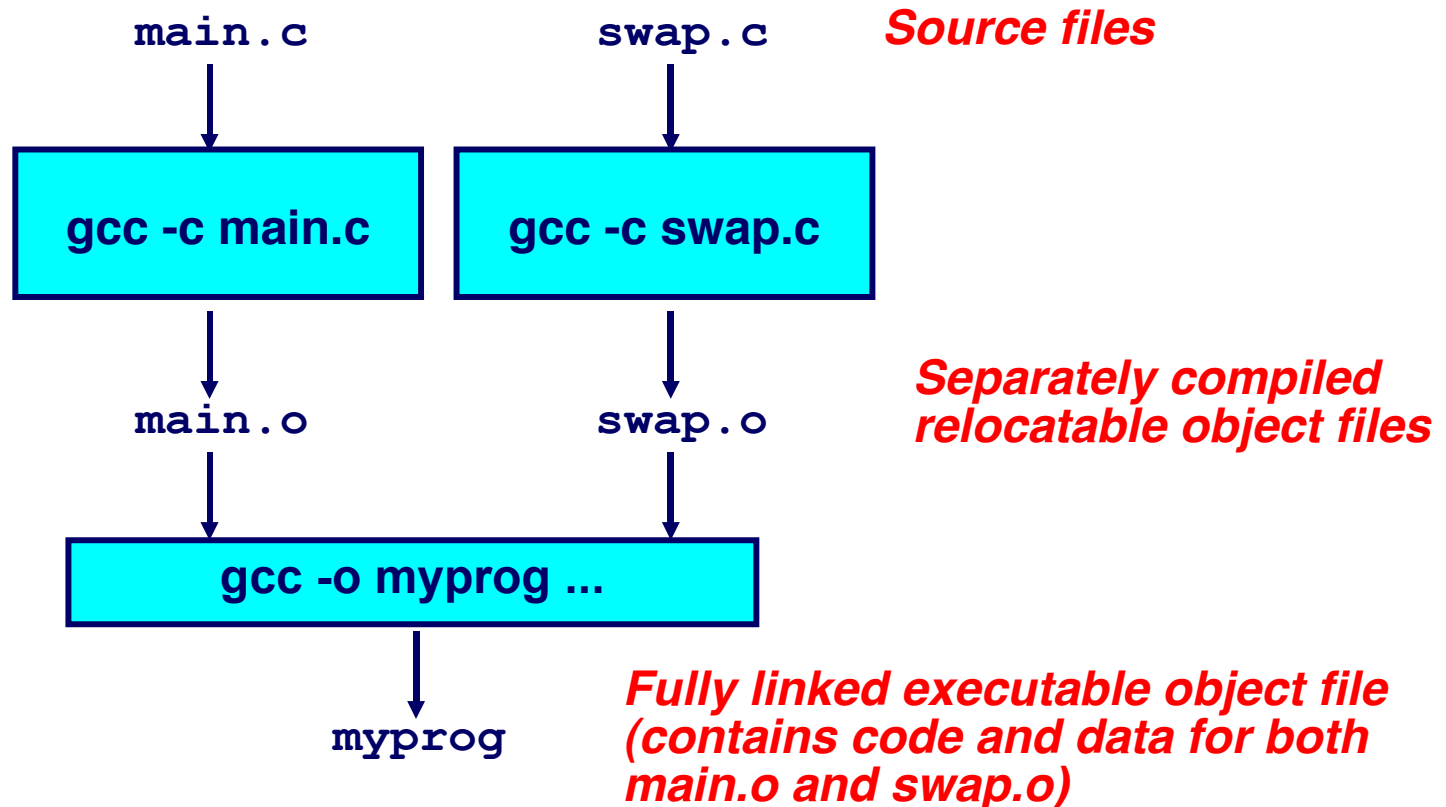
    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
}
```

References
to buf[]

Static Linking

Programs are translated and linked using a *compiler driver*:

- `unix> gcc -c main.c` # Compile .c to .o
- `unix> gcc -c swap.c`
- `unix> gcc -o myprog main.o swap.o` # This runs ld with
the right flags



The linker's problem...

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

void swap()
{
    int temp;

    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
}
```

- Find the definition of each **undefined reference** in an object file
 - Example: The use of “swap()” in main.o needs to be resolved to the definition found in swap.o
- **Resolve** all missing references in the code
 - We can compile main.c to main.o without knowing the memory location of swap()
 - So, main.o will contain a “dummy” address in the call instruction.
 - The linker must patch the assembly code in main.o to reference the correct location of swap() at link time.

Why Linkers?

Reason 1: Modularity

- Program can be written as a collection of smaller source files, rather than one monolithic mass.
- Can build libraries of common functions (more on this later)
 - e.g., Math library, standard C library

Why Linkers? (cont)

Reason 2: Efficiency

- Time efficiency: Separate Compilation
 - Change one source file, recompile just that file, and then relink the executable.
 - No need to recompile other source files.
- Space efficiency: Libraries
 - Common functions can be combined into a single library file...
 - Yet executable files contain only code for the functions they actually use.

Disassembly of main.o

```
int buf[2] = {1,2};
```

```
int main()
{
    swap();
    return 0;
}
```

00000000 <main>:

...

a: 55

b: 89 e5

d: 51

e: 83 ec 04

11: e8 **fc ff ff ff**

16: b8 00 00 00 00

...

push %ebp

mov %esp,%ebp

push %ecx

sub \$0x4,%esp

call **12 <main+0x12>**

mov \$0x0,%eax

This is a bogus address!

Just a placeholder for
“swap” (which we don't know the
address of ... yet!)

main.o object file

main.c

```
int buf[2] = {1,2};

int main()
{
    swap();
    return 0;
}
```

Relocation info tells the linker where references to external symbols are in the code

main.o

.text section

code for main()



.data section

buf[]

symbol table

<u>name</u>	<u>section</u>	<u>off</u>
main	.text	0
buf	.data	0
swap	undefined	

relocation info
for .text

<u>name</u>	<u>offset</u>
swap	12

Disassembly of swap.o

```
extern int buf[];

void swap()
{
    int temp;

    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
}
```

Again, these are placeholders.

0x0 refers to buf[0]
0x4 refers to buf[1]

00000000 <swap>:

```
0: 55
1: 89 e5
3: 8b 15 04 00 00 00
9: a1 00 00 00 00
e: a3 04 00 00 00
13: 89 15 00 00 00 00
19: 5d
1a: c3
```

```
push    %ebp
mov     %esp,%ebp
mov     0x4,%edx
mov     0x0,%eax
mov     %eax,0x4
mov     %edx,0x0
pop     %ebp
ret
```

swap.o object file

swap.o

swap.c


```
extern int buf[];

void swap()
{
    int temp;

    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
}
```

.text section

code for swap()



symbol table

name	section	off
swap	.text	0
buf	undefined	

relocation info
for .text

name	offset
buf	0x5
buf	0xa
buf	0xf
buf	0x15

The linker

- The linker takes multiple object files and combines them into a single executable file.
- Three basic jobs:
 - 1) **Copy** code and data from each object file to the executable
 - 2) **Resolve** references between object files
 - 3) **Relocate** symbols to use absolute memory addresses, rather than relative addresses.

Linker operation

```
00000000 <swap>:  
0: push    %ebp  
1: mov     %esp,%ebp  
...
```

p.o

.text section

code for main()
[redacted]

.data section

buf[]

symbol table

name	section	off
main	.text	0
buf	.data	0
swap	undefined	

relocation info
for .text

name	offset
swap	12

.text section

code for swap()
[redacted]
[redacted]
[redacted]
[redacted]

symbol table

name	section	off
swap	.text	0
buf	undefined	

relocation info
for .text

name	offset
buf	0x5
buf	0xa
buf	0xf
buf	0x15

myprog

.text section

code for main()
[redacted]

code for swap()
[redacted]
[redacted]
[redacted]
[redacted]

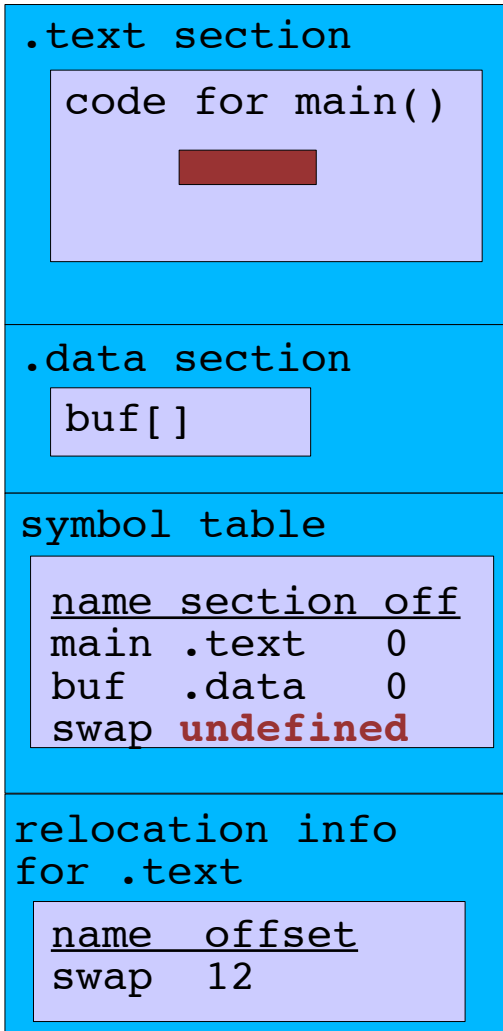
.data section

```
08048368 <swap>:  
8048368: push    %ebp  
8048369: mov     %esp,%ebp  
...
```

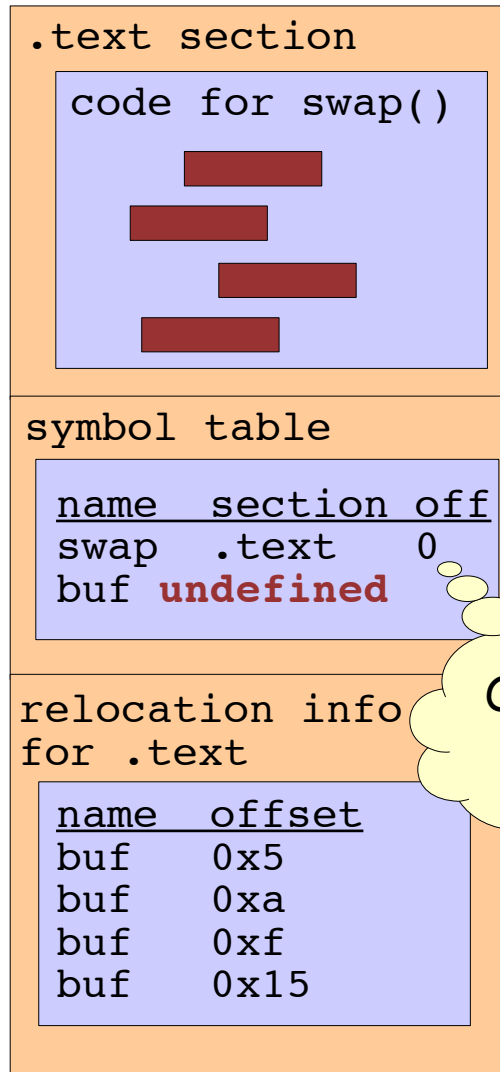
*The executable file contains the **absolute** memory addresses of the code and data (that is, where they will be located when the program actually runs!)*

Linker operation


main.o



swap.o



myprog

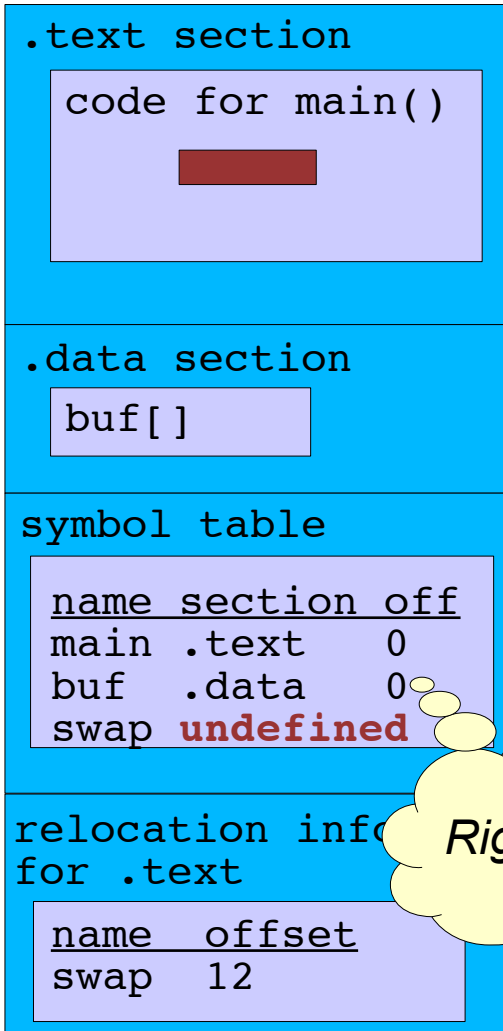


Where's swap()?

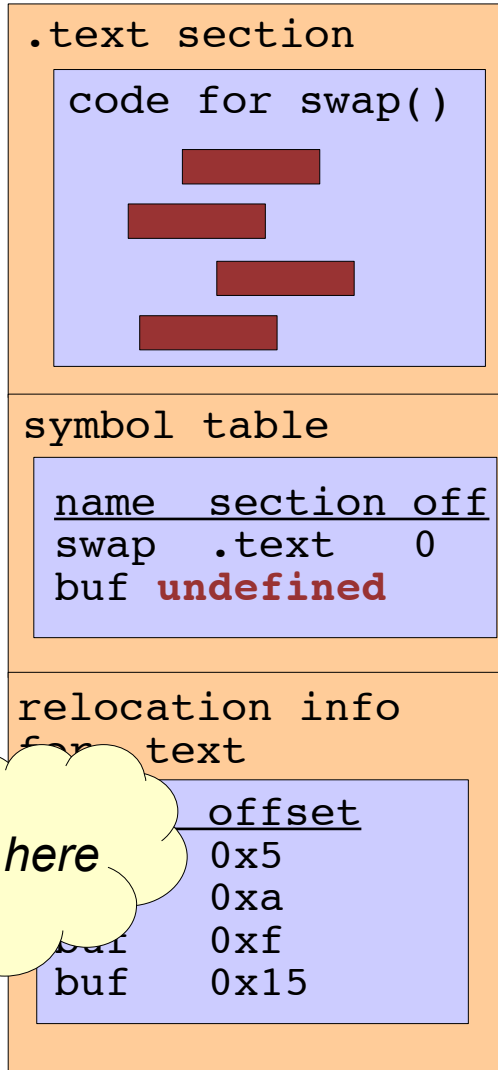
Oh, there it is!

Linker operation

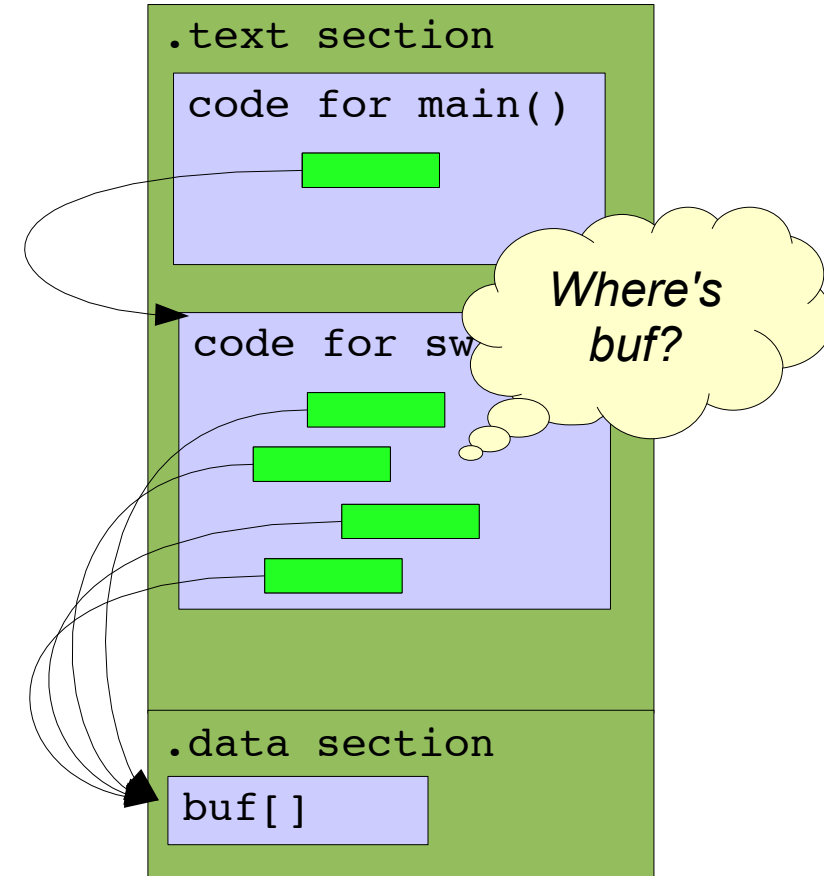
main.o



swap.o



myprog



Right here

Disassembly after linking

08048344 <main>:

What's up with this?
Not the same as
0x8048368...

```

...
804834e: 55          push    %ebp
804834f: 89 e5       mov     %esp, %ebp
8048351: 51          push    %ecx
8048352: 83 ec 04    sub     $0x4, %esp
8048355: e8 0e 00 00 00 call    8048368 <swap>
804835a: b8 00 00 00 00 mov     $0x0, %eax
...

```

08048368 <swap>:

```

8048368: 55          push    %ebp
8048369: 89 e5       mov     %esp, %ebp
804836b: 8b 15 5c 95 04 08 mov     0x804955c, %edx
8048371: a1 58 95 04 08 mov     0x8049558, %eax
8048376: a3 5c 95 04 08 mov     %eax, 0x804955c
804837b: 89 15 58 95 04 08 mov     %edx, 0x8049558
8048381: 5d          pop     %ebp
8048382: c3          ret

```

```
$ objdump -t myprog
```

```

...
08049558 g      0 .data 00000008          buf
...

```

Address	Section	Size	Symbol name
---------	---------	------	-------------

Strong vs. Weak Symbols

- The compiler exports each global symbol as either **strong** or **weak**

- **Strong symbols:**

- Functions
- Initialized global variables



- **Weak symbols:**

- Uninitialized global variables



Strong vs. Weak Symbols

main.c

Strong

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

Strong

Weak

swap.c

Not a global symbol!
(Just a reference.)

```
int myvar;
extern int buf[];
```

Strong

```
void swap()
```

```
{
    int temp;
```

Not a global symbol.

```
    temp = buf[1];
    buf[1] = buf[0];
    buf[0] = temp;
```

```
}
```

Linker rules

- Rule 1: Multiple **strong** symbols with the same name are not allowed.

foo1.c

```
int somefunc() {  
    return 0;  
}
```

foo2.c

```
int somefunc() {  
    return 1;  
}
```

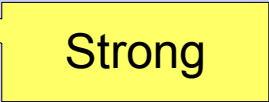
```
$ gcc foo1.c foo2.c  
ld: duplicate symbol _somefunc in /var/folders/Yt/Yta+  
ZMkTGjWQDZZ0edwfCE++UHg/-Tmp-//ccAsmhEJ.o and  
/var/folders/Yt/Yta+ZMkTGjWQDZZ0edwfCE++UHg/-Tmp-//ccVxxcKX.o  
collect2: ld returned 1 exit status
```

Linker rules

- Rule 2: Given a **strong** symbol and multiple **weak** symbols, choose the **strong** symbol.

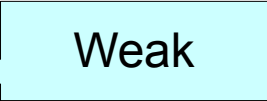
foo1.c

```
void f(void);  
int x = 38;  
  
int main() {  
    f();  
    printf("x = %d\n", x);  
    return 0;  
}
```



foo2.c

```
int x;  
  
void f() {  
    x = 42;  
}
```



```
$ gcc -o myprog foo1.c foo2.c  
$ ./myprog  
x = 42
```

You might not expect this to happen but it does!

Linker rules

- This can lead to some pretty weird bugs!!!

foo1.c

```
void f(void);  
int x = 38;  
int y = 39;  
  
int main() {  
    f();  
    printf("x = %d\n", x);  
    printf("y = %d\n", y);  
    return 0;  
}
```

Strong

foo2.c

```
double x;  
  
void f() {  
    x = 42.0;  
}
```

Weak

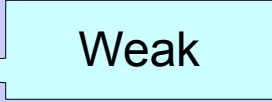
*"double x" is 8 bytes in size!
But resolves to address
of "int x" in foo1.c.*

```
$ gcc -o myprog foo1.c foo2.c  
$ ./myprog  
x = 0  
y = 1078263808
```

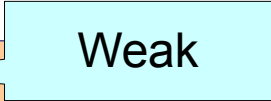
Linker rules

- Rule 3: Given multiple **weak** symbols, pick any one of them

foo1.c

```
void f(void);  
int x;   
  
int main() {  
    x = 38;  
    f();  
    printf("x = %d\n", x);  
    return 0;  
}
```

foo2.c

```
int x;   
  
void f() {  
    x = 42;  
}
```

```
$ gcc -o myprog foo1.c foo2.c  
$ ./myprog  
x = 42
```


Executable and Linkable Format (ELF)

Standard binary format for object files

Originally proposed by AT&T System V Unix

- Later adopted by BSD Unix variants and Linux

One unified format for

- Relocatable object files (.o)
- Shared object files (.so)
- Executable files

ELF Object File Format

Elf header

- Magic number, type (.o, exec, .so), machine, byte ordering, etc.

Segment header table

- Page size, virtual addresses memory segments (sections), segment sizes.

.text section

- Code

.data section

- Initialized global variables

.bss section

- Uninitialized global variables
- “Block Started by Symbol”
- “Better Save Space”
- Has section header but occupies no space

ELF header
Segment header table (required for executables)
.text section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

0

ELF Object File Format (cont)

`.symtab` section

- Symbol table
- Procedure and static variable names
- Section names and locations

`.rel.text` section

- Relocation info for `.text` section
- Addresses of instructions that will need to be modified in the executable
- Instructions for modifying.

`.rel.data` section

- Relocation info for `.data` section
- Addresses of pointer data that will need to be modified in the merged executable

`.debug` section

- Info for symbolic debugging (`gcc -g`)

Section header table

- Offsets and sizes of each section

ELF header
Segment header table (required for executables)
<code>.text</code> section
<code>.data</code> section
<code>.bss</code> section
<code>.symtab</code> section
<code>.rel.text</code> section
<code>.rel.data</code> section
<code>.debug</code> section
Section header table

0

objdump: Looking at ELF files

- Use the “objdump” tool to peek inside of ELF files (.o, .so, and executable files)
- objdump -h: print out list of sections in the file.

```
$ objdump -h myprog
```

```
myprog:      file format elf32-i386
```

```
Sections:
```

Idx	Name	Size	VMA	LMA	File off	Algn
0	.interp	00000013	08048134	08048134	00000134	2**0
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
1	.note.ABI-tag	00000020	08048148	08048148	00000148	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					
2	.hash	00000028	08048168	08048168	00000168	2**2
	CONTENTS, ALLOC, LOAD, READONLY, DATA					

```
...
```

objdump: Looking at ELF files

- Use the “objdump” tool to peek inside of ELF files (.o, .so, and executable files)
- `objdump -s`: print out full contents of each section.

```
$ objdump -s myprog
myprog:          file format elf32-i386

Contents of section .interp:
 8048134 2f6c6962 2f6c642d 6c696e75 782e736f  /lib/ld-linux.so
 8048144 2e3200                                .2.
Contents of section .note.ABI-tag:
 8048148 04000000 10000000 01000000 474e5500 .....GNU.
 8048158 00000000 02000000 06000000 08000000 .....
Contents of section .hash:
 8048168 03000000 05000000 01000000 02000000 .....
 8048178 03000000 00000000 00000000 00000000 .....
```

objdump: Looking at ELF files

- Use the “objdump” tool to peek inside of ELF files (.o, .so, and executable files)
- `objdump -t`: print out contents of symbol table.

```
$ objdump -t myprog
myprog:          file format elf32-i386

SYMBOL TABLE:
08048134 1      d  .interp      00000000      .interp
08048148 1      d  .note.ABI-tag  00000000      .note.ABI-tag
08048168 1      d  .hash        00000000      .hash
08048190 1      d  .gnu.hash    00000000      .gnu.hash
...
```

Packaging Commonly Used Functions

How to package functions commonly used by programmers?

- printf, malloc, strcmp, all that stuff?

Awkward, given the linker framework so far:

- Option 1: Put all functions in a single source file
 - Programmers link big object file into their programs:
 - `gcc -o myprog myprog.o somebighonkinglibraryfile.o`
 - *This would be really inefficient!*
- Option 2: Put each routine in a separate object file
 - Programmers explicitly link appropriate object files into their programs
 - `gcc -o myprog myprog.o printf.o malloc.o free.o strcmp.o strlen.o strerror.o`
 - *More efficient, but a real pain to the programmer*

Static Libraries


Solution: **Static libraries**

- Combine multiple object files into a single **archive** file (file extension “.a”)
- Example: `libc.a` contains a whole slew of object files bundled together.

Linker can also take archive files as input:

- `gcc -o myprog myprog.o /usr/lib/libc.a`
- Linker searches the `.o` files within the `.a` file for needed references and links them into the executable.

Creating Static Libraries



Create a static library file using the UNIX `ar` command

- `ar rs libc.a atoi.o printf.o random.o ...`
- Can list contents of a library using `"ar t libc.a"`

Commonly Used Libraries

`libc.a` (the C standard library)

- 2.8 MB archive of 1400 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

`libm.a` (the C math library)

- 0.5 MB archive of 400 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
...
e_acos.o
e_acosf.o
e_acosh.o
e_acoshf.o
e_acoshl.o
e_acosl.o
e_asin.o
e_asinf.o
e_asinl.o
...
```

Using Static Libraries

Linker's algorithm for resolving external references:

- Scan `.o` files and `.a` files in **command line order**.
- During the scan, keep a list of the current unresolved references.
- As each new `.o` or `.a` file is encountered, try to resolve each unresolved reference in the list against the symbols in the new file.
- If any entries unresolved when done, then return an error.

Problem:

- Command line order matters!
- Moral: put libraries at the end of the command line.

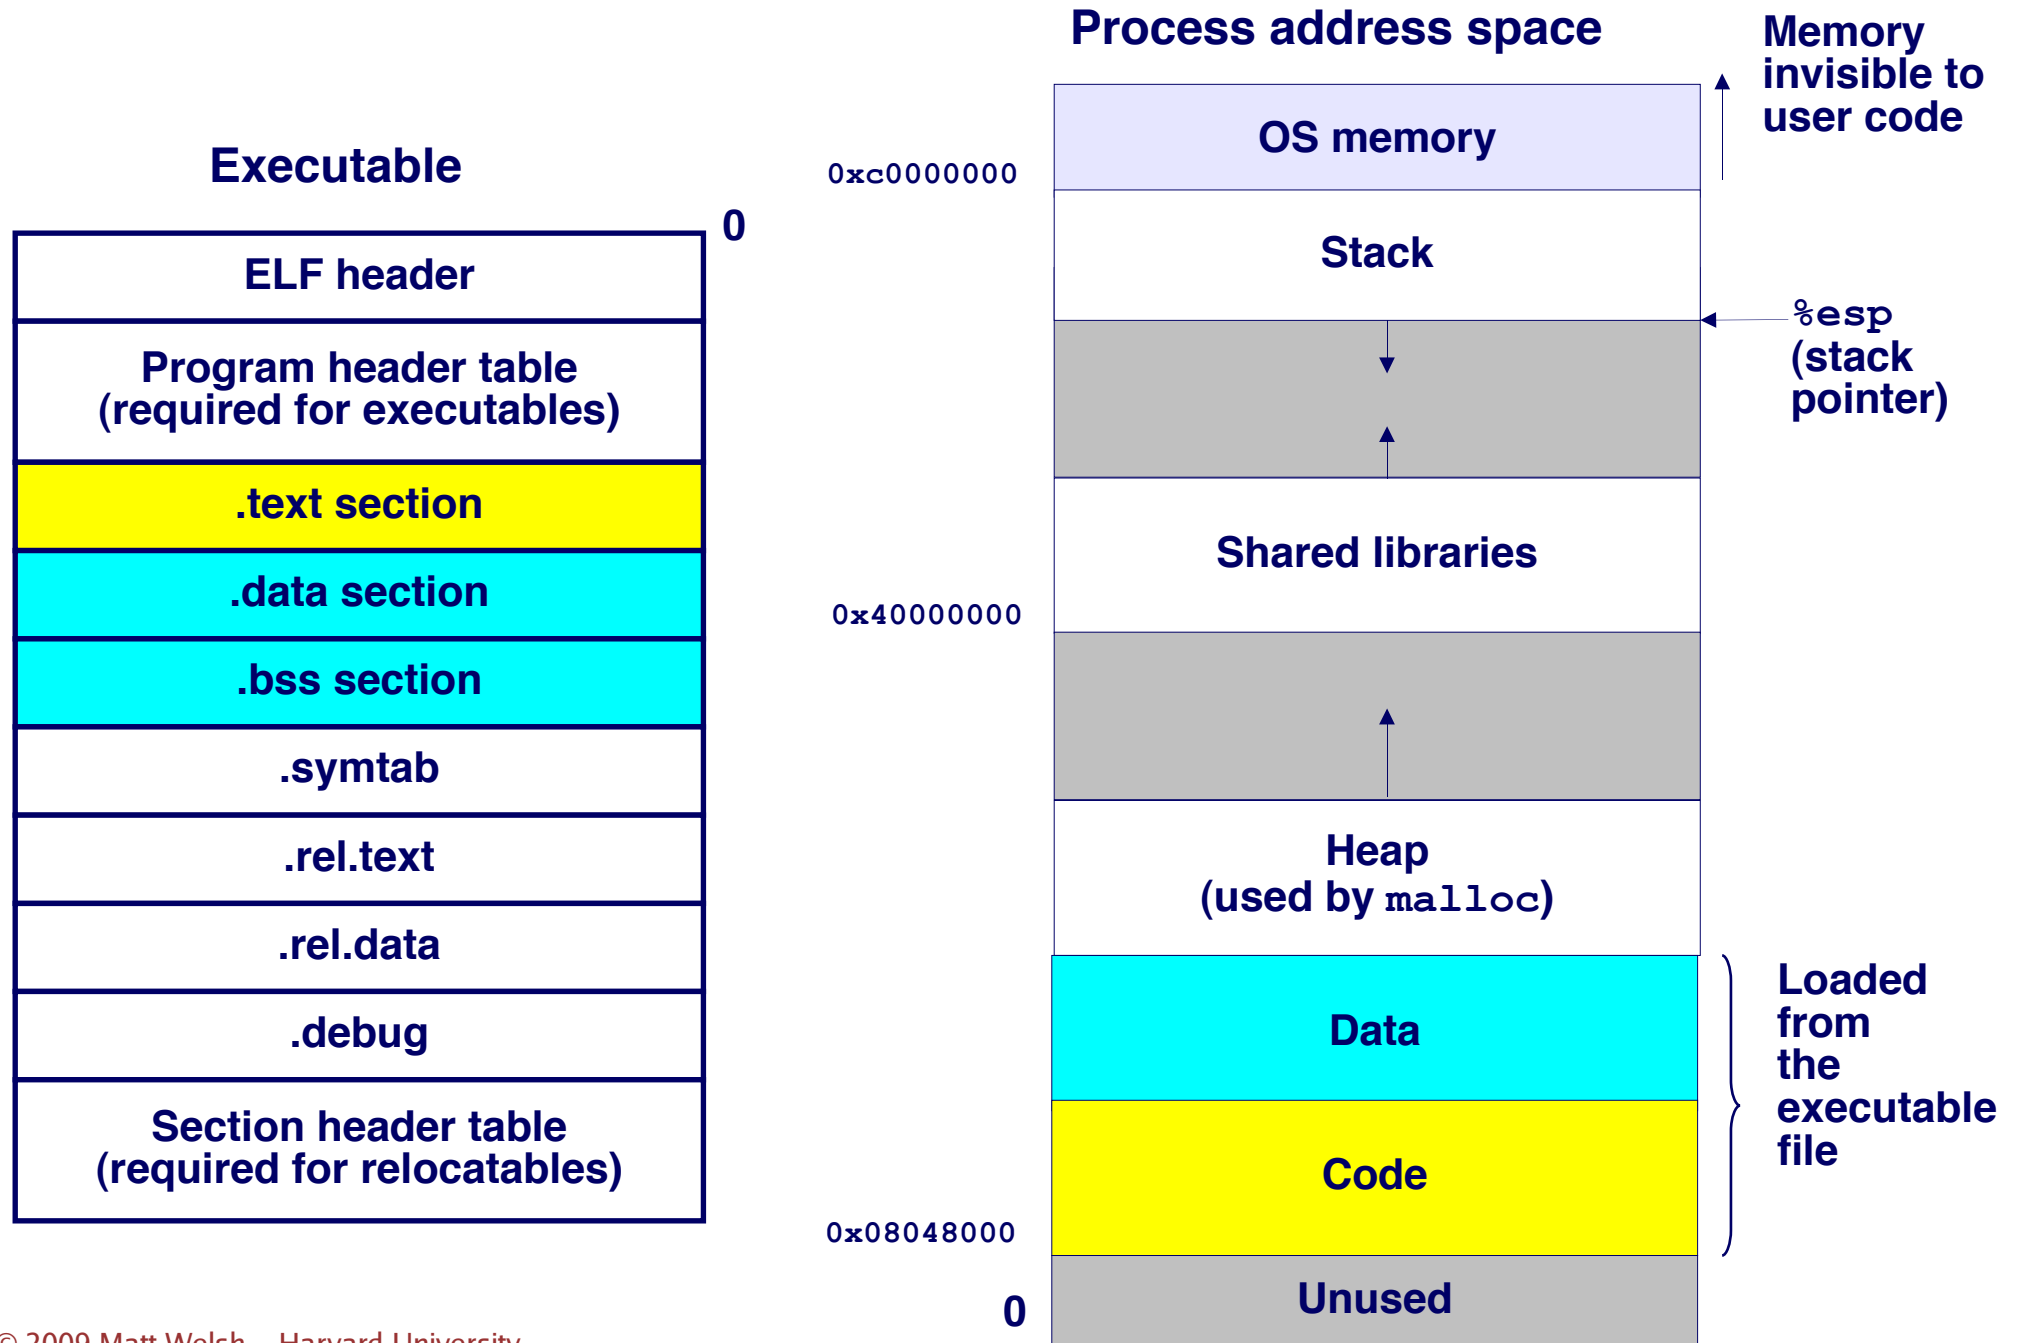
```
unix> gcc mylib.a libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'

unix> gcc libtest.o mylib.a
# works fine!!!
```

Loading an executable

- **Loading** is the process of reading code and data from an executable file and placing it in memory, where it can run.
 - This is done by the operating system.
 - In UNIX, you can use the `execve ()` system call to load and run an executable.
- What happens when the OS runs a program:
 - 1) Create a new **process** to contain the new program (more later this semester!)
 - 2) Allocate memory to hold the program code and data
 - 3) Copy contents of executable file to the newly-allocated memory
 - 4) Jump to the executable's **entry point** (which calls the `main()` function)

Address space layout



Shared Libraries

Static libraries have the following disadvantages:

- Lots of code duplication in the resulting executable files
 - *Every C program needs the standard C library.*
 - *e.g., Every program calling printf() would have a copy of the printf() code in the executable. Very wasteful!*
- Lots of code duplication in the memory image of each running program
 - *OS would have to allocate memory for the standard C library routines being used by every running program!*
- Any changes to system libraries would require relinking every binary!

Solution: **Shared libraries**

- Libraries that are linked into an application *dynamically*, when a program runs!
- On UNIX, “.so” filename extension is used
- On Windows, “.dll” filename extension is used

Shared Libraries

When the OS runs a program, it checks whether the executable was linked against any shared library (.so) files.

- If so, it performs the linking and loading of the shared libraries on the fly.
- Example: `gcc -o myprog main.o /usr/lib/libc.so`
- Can use UNIX `ldd` command to see which shared libs an executable is linked against

```
unix> ldd myprog
```

```
linux-gate.so.1 => (0xfffffe000)  
libc.so.6 => /lib/tls/i686/cmov/libc.so.6 (0xb7df5000)  
/lib/ld-linux.so.2 (0xb7f52000)
```

Can create your own shared libs using `gcc -shared`:

```
gcc -shared -o mylib.so main.o swap.o
```

Runtime dynamic linking

- You can also link against shared libraries at run time!
- Three UNIX routines used for this:
- `dlopen()` -- Open a shared library file
- `dlsym()` -- Look up a symbol in a shared library file
- `dlclose()` -- Close a shared library file
- Why would you ever want to do this?
 - Can load new functionality on the fly, or extend a program after it was originally compiled and linked.
 - Examples include things like browser plug-ins, which the user might download from the Internet well after the original program was compiled and linked.

Dynamic linking example

```
#include <stdio.h>
#include <dlfcn.h>

int main() {
    void *handle;
    void (*somefunc)(int, int);
    char *error;

    /* dynamically load the shared lib that contains somefunc() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
    /* get a pointer to the somefunc() function we just loaded */
    somefunc = dlsym(handle, "somefunc");
    if ((error = dlerror()) != NULL) {
        fprintf(stderr, "%s\n", error);
        exit(1);
    }

    /* Now we can call somefunc() just like any other function */
    somefunc(42, 38);

    /* unload the shared library */
    if (dlclose(handle) < 0) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
    return 0;
}
```

Next time

- Memory and storage technologies
- All about SRAM and DRAM
- How disks work